
RL - Reinforcement Learning


Labs on Policy Evaluation

Prof. Riccardo Berta

2025.10.02

1 Gridworld

Create the environment from the Russell and Norvig's book on AI: **the Gridworld**. This environment is a 3 x 4 grid world in which the agent starts at the bottom-left corner, and it has to reach the top-right corner. There is a hole south of the goal and a wall near the start. The transition function has a 20% noise; that is, 80% the action succeeds, and 20% it fails uniformly at random in orthogonal directions. The reward function is a -0.04 living penalty, a $+1$ for landing on the goal, and a -1 for landing on the hole.



1 - Create the environment Gridworld as described above.

```
# YOUR CODE HERE

# The environment should be a class with the following methods:
# `__init__(self)`: initialize the environment
# `reset(self)`: reset the environment to the initial state
# `step(self, action)`: take an action in the environment.
# The action should be an integer between 0 and 3, where 0 means left,
# 1 means down, 2 means right, and 3 means up.
# The method should return a tuple (state, reward, done),
# where state is the new state, reward is the reward obtained, and done is
# a boolean indicating whether the episode is finished.

UP, DOWN, LEFT, RIGHT = range(4)

class GridWorld:

    # 0:UP, 1:DOWN, 2:LEFT 3:RIGHT
```

```
def __init__(self):
    self.reset()

def reset(self):
    self.observation_space = 12
    self.state = 8
    self.columns = 4
    self.rows = 3
    self.terminated = False
    self.reward = 0
    return self.state

def step(self, action):
    if self.terminated: raise ValueError('Episode has terminated')
    if action not in [UP, DOWN, LEFT, RIGHT]: raise ValueError('Invalid action')

    # add stochasticity
    if action == UP: action = np.random.choice([UP, LEFT, RIGHT], p=[0.8, 0.1,
        ↪ 0.1])
    if action == DOWN: action = np.random.choice([DOWN, LEFT, RIGHT], p=[0.8, 0.1,
        ↪ 0.1])
    if action == LEFT: action = np.random.choice([LEFT, UP, DOWN], p=[0.8, 0.1,
        ↪ 0.1])
    if action == RIGHT: action = np.random.choice([RIGHT, UP, DOWN], p=[0.8, 0.1,
        ↪ 0.1])

    # manage boundaries
    row = self.state // self.columns
    column = self.state % self.columns
    if action == LEFT: column = max(column - 1, 0)
    if action == DOWN: row = min(row + 1, self.rows - 1)
    if action == RIGHT: column = min(column + 1, self.columns - 1)
    if action == UP: row = max(row - 1, 0)

    # manage obstacles
    old_state = self.state
    self.state = row * self.columns + column
    if self.state == 5: self.state = old_state

    # manage rewards
    self.reward = -0.04
    if self.state == 3: self.terminated = True; self.reward = 1
    elif self.state == 7: self.terminated = True; self.reward = -1
```


```

    return self.state, self.reward, self.terminated, 0, 0

grid_world = GridWorld()

```

2 - Write the following policy for this environment



```

# YOUR CODE HERE

# The policy should be a function that takes an integer between 0 and 11
# (the state) and returns an integer between 0 and 3 (the action).

pi = lambda s: {
    0:RIGHT, 1:RIGHT, 2:RIGHT, 3:LEFT,
    4:UP,    5:LEFT, 6:UP,    7:LEFT,
    8:UP,    9:LEFT, 10:LEFT, 11:LEFT
}[s]

```

3 - Show the policy by printing the selected action for each state:

```

# YOUR CODE HERE

# The function should print the policy in a human-readable format.
# For example, the output could be:
# state: 0 → action: RIGHT
# state: 1 → action: RIGHT
# etc ...

def print_policy(pi, env):
    for state in range(env.observation_space):

```

```
    action = pi(state)
    action_icon = ' '
    if action == UP: action_icon = 'UP'
    if action == DOWN: action_icon = 'DOWN'
    if action == LEFT: action_icon = 'LEFT'
    if action == RIGHT: action_icon = 'RIGHT'
    print('state:', state, '→', 'action:', action_icon)

print_policy(pi, grid_world)
```

```
state: 0 -> action: RIGHT
state: 1 -> action: RIGHT
state: 2 -> action: RIGHT
state: 3 -> action: LEFT
state: 4 -> action: UP
state: 5 -> action: LEFT
state: 6 -> action: UP
state: 7 -> action: LEFT
state: 8 -> action: UP
state: 9 -> action: LEFT
state: 10 -> action: LEFT
state: 11 -> action: LEFT
```

4 - Now, evaluate this policy using $TD(\lambda)$

```
# YOUR CODE HERE

# you need to decay alpha

import numpy as np

def decay_alpha(init_value, min_value, decay_steps, max_steps):

    # calculate the number of the remaining steps after the decay
    rem_steps = max_steps - decay_steps

    # logspace returns numbers spaced evenly on a log scale
    # base^start is the starting value of the sequence,
    # base^stop is the final value of the sequence
    # num is the number of values to generate
```

```
# base is the base of the log space
values = np.logspace(start=0, stop=-2, num=decay_steps, base=10)

# because the values may not end exactly at 0, given it's the log,
# we change them to be between 0 and 1 so that the curve looks smooth and nice.
values = (values - values.min()) / (values.max() - values.min())

# linear transformation and get points between init_value and min_value
values = (init_value - min_value) * values + min_value

# repeats the rightmost value rem_step number of times
values = np.pad(values, (0, rem_steps), 'edge')

return values
```

```
# YOUR CODE HERE

# you need to decay the discount

def decay_discounts(gamma, max_steps):
    discounts = np.logspace(start=0, stop=max_steps, num=max_steps, base=gamma,
↪ endpoint=False);
    return discounts
```

```
# YOUR CODE HERE

# you need the td_lambda implementation

def td_lambda(pi, env, gamma=1.0, lambda_=0.2,
              init_alpha=0.5, min_alpha=0.01,
              decay_episodes=350, n_episodes=500):

    # calculate all alphas at once
    alphas = decay_alpha(init_alpha, min_alpha, decay_episodes, n_episodes);

    # initialize the current estimate of the state-value function
    v = np.zeros(env.observation_space, dtype=float)

    # create a list to save copies of v for offline analysis
    v_track = np.zeros((n_episodes, env.observation_space), dtype=float)
```

```
# initialize the eligibility trace vector
E = np.zeros(env.observation_space, dtype=float)

# loop for every episode
for e in range(n_episodes):

    # set E to zero every new episode
    E.fill(0)

    # get the initial state
    state, done = env.reset(), False

    # get into the time step loop
    while not done:

        # sample the policy pi for the action to take in state
        action = pi(state)

        # interact with the environment for one step and get the experience tuple
        next_state, reward, terminated, truncated, info = env.step(action)
        if(terminated or truncated):
            done = True;

        # use that experience to calculate the TD error as usual
        td_target = reward + gamma * v[next_state] * (not done)
        td_error = td_target - v[state]

        # increment the eligibility of state by 1
        E[state] = E[state] + 1

        # apply the error update to all eligible states as indicated by E
        v = v + alphas[e] * td_error * E

        # decay E
        E = lambda_ * E

        # update the state variable for the next iteration
        state = next_state

    v_track[e] = v

return v, v_track
```

```
# YOUR CODE HERE

# run the algorithm and print the value function

v_td_lambda, v_td_lambda_track = td_lambda(pi, grid_world, lambda_=0.3,
    ↪ n_episodes=500)
print(v_td_lambda)
```

```
[0.8415656  0.89934748 0.95800912 0.          0.79107068 0.
 0.61815666 0.          0.74072074 0.60860885 0.          0.          ]
```

5 - Plot the estimated value for state 8, 6 and 2 for all episodes

```
# YOUR CODE HERE

# plot the value function over episodes for state 8, 6 and 2

import matplotlib.pyplot as plt

plt.figure(figsize=(10,6))
legends = ['v(8)', 'v(6)', 'v(2)']
plt.plot(v_td_lambda_track[:,8])
plt.plot(v_td_lambda_track[:,6])
plt.plot(v_td_lambda_track[:,2])
plt.title('TD(lambda) estimates over episodes')
plt.ylabel('State-Value function')
plt.xlabel('Episodes')
plt.legend(legends)

plt.show()
```